

Chapter 3

SQL



- Background
- Basic Structure__query
- DDL
- DML(Insert ,Delete,Update)
- Embedded SQL
- Dynamic SQL
- ODBC,JDBC
- Transact_SQL



Structured query language (SQL)

1970s , IBM Sequel(System R)

1986, ANSI and ISO , **SQL_86** ;

1987, IBM, SAA_SQL;

1989, ANSI, SQL_89, an extended standard for SQL;

SQL_92;

1999 SQL;

SQL:2003



- Data definition language (DDL);

- Data manipulation language (DML);

- Data control language (DCL);

 - Integrity;

 - Authorization;

SQL is based on set and relational operations with certain modifications and enhancements

- Embedded SQL and dynamic SQL;



- A typical SQL query has the form:

***select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P***

A_i represent attributes

r_i represent relations

P is a predicate.

- This query is equivalent to the relational algebra expression.

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

The result of an SQL query is a relation



- The **select** clause corresponds to the **projection** operation of the relational algebra. It is used to list the attributes desired in the result of a query.

- Find the names of all instructors:

select name
from instructor

- In the “pure” relational algebra syntax, the query would be:

$\Pi_{\text{name}}(\text{instructor})$

- An asterisk in the select clause denotes “all attributes”

select *
from instructor



- The **select** clause can contain arithmetic expressions involving the operation, +, −, *, and /, and operating on constants or attributes of tuples.

The query:

```
select ID, name, salary/12  
from instructor
```

would return a relation that is the same as the instructor relation, except that the value of the attribute salary is divided by 12.

Can rename “salary/12” using the as clause:

```
select ID, name, salary/12 as monthly_salary
```



- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after **select**.
- **Find the department names of all instructors, and remove duplicates**

***select distinct dept_name
from instructor***

- The keyword **all** specifies that duplicates should not be removed.

***select all dept_name
from instructor***



- The **where** clause corresponds to the **selection** predicate of the relational algebra.

- Find all instructors in Comp. Sci. dept

select name
from instructor
where dept_name = 'Comp. Sci.'



- Comparison results can be combined using the logical connectives and, or, and not
 - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and  
salary > 80000
```

Comparisons can be applied to results of arithmetic expressions.

Between..and



- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000
 - **select** *name*
from *instructor*
where *salary* **between** 90000 **and** 100000
- Tuple comparison
 - **select** *name, course_id*
from *instructor, teaches*
where (*instructor.ID, dept_name*) = (*teaches.ID, 'Biology'*);



The from Clause

- The **from** clause corresponds to the Cartesian product operation of the relational algebra. It lists the relations to be scanned in the evaluation of the expression.

- Find the Cartesian product instructor X teaches

select *

from instructor, teaches

For common attributes (e.g., ID), the attributes in the resulting table are renamed using the relation name (e.g., instructor.ID)

- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).



- *Find the names of all instructors who have taught some course and the course_id*
select name, course_id
from instructor , teaches
where instructor.ID = teaches.ID
- *Find the names of all instructors in the Art department who have taught some course and the course_id*
select name, course_id
from instructor , teaches
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'



The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

old_name as new_name

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

select distinct T.name

from instructor as T, instructor as S

where T.salary > S.salary and S.dept_name = 'Comp. Sci.'

Keyword as is optional and may be omitted

instructor as T \equiv instructor T

Alias



- percent (%). The % character matches any substring;
- underscore (_). The _ character matches any character;
- Find the names of all instructors whose name includes the substring “dar”.

```
select name  
from instructor  
where name like '%dar%'
```



■ Match the name “Main%”

like 'Main\%' escape '\'

■ SQL supports a variety of string operations such as

- concatenation (using “||”)
- converting from upper to lower case (and vice versa)
- finding string length, extracting substrings, etc.



● Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

select distinct *name*
from *instructor*
order by *name*

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

Example: **order by** *name* **desc**

Can sort on multiple attributes

Example: **order by** *dept_name, name*



- In relations with duplicates, SQL can define how many copies of tuples appear in the result.
- **Multiset** versions of some of the relational algebra operators – given multiset relations r_1 and r_2 :
 1. $\sigma_{\theta}(r_1)$: If there are c_1 copies of tuple t_1 in r_1 , and t_1 satisfies selections σ_{θ} , then there are c_1 copies of t_1 in $\sigma_{\theta}(r_1)$.
 2. $\Pi_A(r)$: For each copy of tuple t_l in r_1 , there is a copy of tuple $\Pi_A(t_l)$ in $\Pi_A(r_1)$ where $\Pi_A(t_l)$ denotes the projection of the single tuple t_l .
 3. $r_1 \times r_2$: If there are c_1 copies of tuple t_1 in r_1 and c_2 copies of tuple t_2 in r_2 , there are $c_1 \times c_2$ copies of the tuple t_1, t_2 in $r_1 \times r_2$



- Example: Suppose multiset relations $r_1 (A, B)$ and $r_2 (C)$ are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then $\Pi_B(r_1)$ would be $\{(a), (a)\}$, while $\Pi_B(r_1) \times r_2$ would be

$$\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$$

- SQL duplicate semantics:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

is equivalent to the *multiset* version of the expression:



- The set operations **union**, **intersect**, and **except** operate on relations and correspond to the relational algebra operations \cup , \cap , $-$.
- Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.



Set Operations_Example

- Find courses that ran in Fall 2009 or in Spring 2010
(select course_id from section where sem = 'Fall' and year = 2009)
union
(select course_id from section where sem = 'Spring' and year = 2010)
- Find courses that ran in Fall 2009 and in Spring 2010
(select course_id from section where sem = 'Fall' and year = 2009)
intersect
(select course_id from section where sem = 'Spring' and year = 2010)
- Find courses that ran in Fall 2009 but not in Spring 2010
(select course_id from section where sem = 'Fall' and year = 2009)
except
(select course_id from section where sem = 'Spring' and year = 2010)



These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values



Aggregate Functions_Example *DataBase System Concepts*

- Find the average salary of instructors in the Computer Science department

- **select avg** (*salary*)
from *instructor*
where *dept_name*= 'Comp. Sci.';

- Find the total number of instructors who teach a course in the Spring 2010 semester

- **select count** (**distinct** *ID*)
from *teaches*
where *semester* = 'Spring' **and** *year* = 2010;

- Find the number of tuples in the *course* relation

- **select count** (*)
from *course*;



- Find the average salary of instructors in each department

- **select** dept_name, avg (salary) as avg_salary
from instructor
group by dept_name;

Note: Attributes in **select** clause outside of aggregate functions must appear in **group by** list

- /* erroneous query */
select dept_name, ID, avg (salary)
from instructor
group by dept_name;



Aggregate Functions – Group By *DataBase System Concepts*

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 76766 | Crick | Biology | 72000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 12121 | Wu | Finance | 90000 |
| 76543 | Singh | Finance | 80000 |
| 32343 | El Said | History | 60000 |
| 58583 | Califieri | History | 62000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 22222 | Einstein | Physics | 95000 |

| <i>dept_name</i> | <i>salary</i> |
|------------------|---------------|
| Biology | 72000 |
| Comp. Sci. | 77333 |
| Elec. Eng. | 80000 |
| Finance | 85000 |
| History | 61000 |
| Music | 40000 |
| Physics | 91000 |



Aggregate Functions_Having clause *DataBase System Concepts*

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups



- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
 - Example: $5 + \text{null}$ returns null
- The predicate **is null** can be used to check for null values.
 - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



- Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
 - count returns 0
 - all other aggregates return null



- Any comparison with *null* returns *unknown*
 - Example: $5 < \text{null}$ or $\text{null} <> \text{null}$ or $\text{null} = \text{null}$
- Three-valued logic using the truth value *unknown*:
 - OR: $(\text{unknown} \textbf{ or } \text{true}) = \text{true}$,
 $(\text{unknown} \textbf{ or } \text{false}) = \text{unknown}$
 $(\text{unknown} \textbf{ or } \text{unknown}) = \text{unknown}$
 - AND: $(\text{true} \textbf{ and } \text{unknown}) = \text{unknown}$,
 $(\text{false} \textbf{ and } \text{unknown}) = \text{false}$,
 $(\text{unknown} \textbf{ and } \text{unknown}) = \text{unknown}$
 - NOT: $(\textbf{not } \text{unknown}) = \text{unknown}$
 - “*P* is **unknown**” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*



Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.

- The nesting can be done in the following SQL query

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

as follows:

- A_i can be replaced by a subquery that generates a single value.
- r_i can be replaced by any valid subquery
- P can be replaced with an expression of the form:

B <operation> (subquery)

Where B is an attribute and <operation> to be defined later.



Subqueries in the Where Clause



● ***Subqueries in the Where Clause*** *DataBase System Concepts*

- • A common use of subqueries is to perform tests:
 - – For set membership
 - – For set comparisons
 - – For set cardinality.
-
-
-
-
-



- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id  
from section  
where semester = 'Fall' and year= 2009 and  
       course_id in (select course_id  
                     from section  
                     where semester = 'Spring'  
                     and year= 2010);
```



- Find courses offered in Fall 2009 but not in Spring 2010

*select distinct course_id
from section*

*where semester = 'Fall' and year= 2009 and
course_id not in (select course_id
from section*

*where semester = 'Spring' and
year= 2010);*



Set Membership (Cont.)

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

*select count (distinct ID)
from takes*

*where (course_id, sec_id, semester, year) in
(select course_id, sec_id, semester, year
from teaches
where teaches.ID= 10101);*

■ ***Note: Above query can be written in a much simpler manner.***

The formulation above is simply to illustrate SQL features.



Set Comparison – “some” Clause *DataBase System Concepts*

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept  
name = 'Biology';
```

- Same query using > some clause

```
select name  
from instructor  
where salary > some (select salary  
from instructor  
where dept name = 'Biology')
```



Definition of Some Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$
Where $<\text{comp}>$ can be: $<, \leq, >, =, \neq$

$$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$$

$$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$$

$$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$$

$(= \text{some}) \equiv \text{in}$

However, $(\neq \text{some}) \not\equiv \text{not in}$



● **Set Comparison – “all” Clause**

- • Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

*select name
from instructor
where salary > all (select salary
from instructor
where dept name = 'Biology');*



Definition of All Clause

- $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 \text{ < all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 \text{ < all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{ all}) \equiv \text{not in}$

However, $(= \text{ all}) \not\equiv \text{in}$



Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists** $r \Leftrightarrow r \neq \emptyset$
- **not exists** $r \Leftrightarrow r = \emptyset$



Use of “exists” Clause

- Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year =
2010
              and S.course_id = T.course_id);
```

- Correlation name** – variable S in the outer query
- Correlated subquery** – the inner query



Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S  
where not exists ( (select course_id  
                    from course  
                    where dept_name = 'Biology')  
except  
                  (select T.course_id  
                    from takes as T  
                    where S.ID = T.ID));
```

■ **Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$**

■ **Note: Cannot write this query using = all and its variants**



Test for Absence of Duplicate Tuples

DataBase System Concepts

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates .
- Find all courses that were offered at most once in 2009

```
select T.course_id  
from course as T  
where unique (select R.course_id  
                from section as R  
                where T.course_id= R.course_id  
                    and R.year = 2009);
```



Subqueries in the From Clause



Subqueries in the From Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as
      avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```



Subqueries in the From Clause

DataBase System Concepts

- Note that we do not need to use the having clause
- Another way to write above query

```
select dept_name, avg_salary  
from (select dept_name, avg (salary)  
      from instructor  
      group by dept_name) as dept_avg  
(dept_name, avg_salary)  
where avg_salary > 42000;
```



- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```



Complex Queries using With Clause

DataBase System Concepts

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
        from instructor  
        group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
        from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value > dept_total_avg.value;
```



Subqueries in the Select Clause



Scalar Subquery

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department

```
select dept_name,  
        (select count(*)  
         from instructor  
         where department.dept_name =  
instructor.dept_name)  
        as num_instructors  
from department;
```

- Runtime error if subquery returns more than one result tuple



Select customer_name

From borrower

Where **exists** (select *

from depositor

where depositor.customer_name=borrower.customer_name)

Set operation:

in, not in, >some/all, <=some/all, =some/all, <>some/all,

exists, not exists,

unique, not unique



SELECT [DISTINCT] <Attributes list>

FROM <relations list>

[***WHERE*** < predicate >]

[***GROUP BY*** < Attributes list >

[***HAVING*** < Attributes list >]]

[***ORDER BY*** < Attribute > [ASC|DESC], ..., <
Attribute > [ASC|DESC]];



Data Definition Language (DDL) *DataBase System Concepts*

Allows the specification of not only a set of relations but also information about each relation, including:

- The schema for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- The set of indices to be maintained for each relations.
- Security and authorization information for each relation.
- The physical storage structure of each relation on disk



Create Table Construct

```
create table r (A1 D1, A2 D2, ..., An Dn,  
                (integrity_constraint1),  
                ...,  
                (integrity_constraintk))
```

r is the name of the relation

each *A*_{*i*} is an attribute name in the schema of relation *r*

*D*_{*i*} is the data type of values in the domain of attribute *A*_{*i*}

Example:

```
■ create table instructor (  
    ID           char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary     numeric(8,2))
```



- **not null**
- **primary key** (A_1, \dots, A_n)
- *Foreign Key* (A_1, \dots, A_n) *reference table r on delete restrict/cascade/set NULL*
- **check** (P), where P is a predicate

```
create table instructor (  
    ID          char(5),  
    name        varchar(20) not null,  
    dept_name    varchar(20),  
    salary       numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references  
    department);
```



And a Few More Relation Definitions *DataBase System Concepts*

- **create table** *student* (
 ID **varchar**(5),
 name **varchar**(20) not null,
 dept_name **varchar**(20),
 tot_cred **numeric**(3,0),
 primary key (*ID*),
 foreign key (*dept_name*) **references** *department*);

- **create table** *takes* (
 ID **varchar**(5),
 course_id **varchar**(8),
 sec_id **varchar**(8),
 semester **varchar**(6),
 year **numeric**(4,0),
 grade **varchar**(2),
 primary key (*ID*, *course_id*, *sec_id*, *semester*, *year*) ,
 foreign key (*ID*) **references** *student*,
 foreign key (*course_id*, *sec_id*, *semester*, *year*) **references** *section*);

- Note: *sec_id* can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester



```
■ create table course (  
    course_id    varchar(8),  
    title        varchar(50),  
    dept_name    varchar(20),  
    credits      numeric(2,0),  
    primary key (course_id),  
    foreign key (dept_name) references  
    department);
```



- **char(n).** Fixed length character string, with user_specified length n .
- **varchar(n).** Variable length character strings, with user_specified maximum length n .
- **int.** Integer (a finite subset of the integers that is machine_dependent).
- **smallint.** Small integer (a machine_dependent subset of the integer domain type).
- **numeric(p,d).** Fixed point number, with user_specified precision of p digits, with n digits to the right of decimal point.
- **real, double precision.** Floating point and double_precision floating point numbers, with machine_dependent precision.
- **float(n).** Floating point number, with user_specified precision of at least n digits.



Drop and Alter Table Constructs *DataBase System Concepts*

- The **drop table** command deletes all information about the dropped relation from the database.
- The **alter table** command is used to add attributes to an existing relation.

alter table r add A D

where A is the name of the attribute to be added to relation r and D is the domain of A .



- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a *view*.
- A view provides a mechanism to hide certain data from the view of certain users.
- the command:
create view v as <query expression>
where:
 - 👉 <query expression> is any legal expression
 - 👉 The view name is represented by v



■ Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.

■ When a view is created, the query expression is stored in the database; the expression is substituted into queries using the view.



- A view consisting of branches and their customers

```
create view all_customer as  
  (select branch_name, customer_name  
    from depositor, account  
    where depositor.account_number =  
    account.account_number)  
  union  
  (select branch_name, customer_name  
    from borrower, loan  
    where borrower.loan_number = loan.loan_number)
```

Find all customers of the Perryridge branch

```
select customer_name  
from all_customer  
where branch_name = 'Perryridge'
```



- One view may be used in the expression defining another view. A way to define the meaning of views defined in terms of other views.

- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.

- View expansion of an expression repeats the following replacement step:

- repeat**

- Find any view relation v_i in e_1

- Replace the view relation v_i by the expression defining v_i

- until** no more view relations are present in e_1

- As long as the view definitions are not recursive, this loop will terminate



Modification of the Database

-Deletion

- The command :

delete from

- Delete all instructors

delete from *instructor*

- Delete all instructors from the Finance department

delete from *instructor*

where *dept_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

delete from *instructor*

where *dept name* in (**select** *dept name*

from *department*

where *building* = 'Watson');



- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor  
where salary < (select avg (salary)  
               from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
 1. First, compute avg (salary) and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing avg or retesting the tuples)



Modification of the Database -Insertion

- The command :

Insert into Values()

- Add a new tuple to *course*

insert into course

***values ('CS-437', 'Database Systems', 'Comp. Sci.',
4);***

insert into course (course_id, title, dept_name, credits)

***values ('CS-437', 'Database Systems', 'Comp. Sci.',
4);***

- Add a new tuple to *student* with *tot_creds* set to null

insert into student

values ('3003', 'Green', 'Finance', null);



- Add all instructors to the *student* relation with tot_creds set to 0

```
insert into student  
select ID, name, dept_name, 0  
from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



Modification of the Database -Update

- The command :

Update Set

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

update instructor

*set salary = salary * 1.03*

where salary > 100000;

update instructor

*set salary = salary * 1.05*

where salary <= 100000;

- The order is important
- Can be done better using the case statement



Case Statement for Conditional Updates

- Same query as before but with case statement

update instructor

set salary = case

when salary <= 100000 then salary * 1.05

else salary * 1.03

end



Updates with Scalar Subqueries *DataBase System Concepts*

- Recompute and update tot_creds value for all students

```
update student S  
  set tot_cred = (select sum(credits)  
                  from takes, course  
                  where takes.course_id =  
course.course_id and  
                  S.ID= takes.ID.and  
                  takes.grade <> 'F' and  
                  takes.grade is not null);
```

- Sets tot_creds to null for students who have not taken any course
- Instead of **sum(credits)**, use:

```
  case  
    when sum(credits) is not null then sum(credits)  
    else 0  
  end
```



CREATE INDEX H_INDEX ON STUDENT(HEIGHT)

**CREATE UNIQUE INDEX SC_INDEX ON
SC(SNO ASC,CNO DESC)**

DROP INDEX H_INDEX



Query

SELECT..

Modification

Delete, Insert , update

Definition

Table, View

